

Deadlock: Occur with multiple process when an requires a resource held by others mutually. Deadlock can be eliminated between n processes **accessing resource than n .**
Pthread has `pthread_mutex_trylock()` to prevent deadlocks.

Semaphores
A semaphore, s , is a positive integer (including 0) operated upon by two operations:
P Operation: waits until s is greater than 0 and then decrements s by 1 and allows process to continue.
V Operation: increments s by 1 to release one of the waiting processes (if any)
Notice: these operations are performed indivisibly. s can take on positive values other than zero and one, meaning of recording the number of "resource unites" available or used and can be used to solve producer/consumer problems. Semaphore routines exist for **UNIX processes**, though they can be written and they do exist in real-time extension to Pthreads.

Monitor
A suit of procedures that provides the only method to access a shared resource. Reading and writing can only be done by using a monitor procedure, and only one process can use a monitor at any instant.

Condition Variables
A critical section is to be executed if a specific global condition exist. With locks the global variable would need to be examined at frequent intervals within a critical section, which is time-consuming. This can be solve by **condition variable**

Bernstein's Condition
Condition that sufficient to determine whether two processes can be executed simultaneously:

I_1 is the set of memory locations read by process P_1
 O_1 is the set of memory locations altered by process P_1
For two process P_1 and P_2 to be execute simultaneously:
 $I_1 \cap O_2 = \emptyset$
 $I_2 \cap O_1 = \emptyset$
 $O_1 \cap O_2 = \emptyset$
The set of outputs of each process must also be different:

Shared Data in Systems with Caches
Cache Coherence Protocols: In the **update policy**, copies of data in all caches are updated at the time one copy is altered. In the **invalidate policy**, when one copy of data is altered, the same data in any other cache is **invalidated**. These copies are only updated when the associated processor makes reference for it.
False Sharing: Different parts of block required by different processors but not same bytes. If one processor writes to one part of the block, copies of the complete block in other caches must be updated or invalidated though the actual data is not shared.
Compiler to alter the layout of the data stored in the main memory, separating data only altered by one processor into different blocks.

Quinn 17
Shared-memory Model: #active threads 1 at start and finish of the program, changes dynamically during execution. Execute and profile sequential program. Incrementally make it parallel. Stop when further effort not warranted.

Message-passing Model: all processes active through execution. Sequential-to-parallel transformation requires major effort. Transformation done in one giant step.

Incremental Parallelization:
Sequential program a special case of a share-memory parallel program. Parallel share-memory programs may only have a single parallel loop. **Incremental Parallel:** process of converting a sequential program to a parallel program a little bit at a time.

Execution Context: address space containing all of the variables a thread may access, includes: static variables, dynamically allocated data structures in the heap, variables on run-time stack, additional run-time stack for funcs invoked by thread.

Shared/private variables: same address in threads or not.
Omp_get_num_procs: return #physical processors
Private clause: direct compiler to make some variables private
Firstprivate clause: create private variables having initial values identical to the variable controlled by the master thread as the loop is entered.

Critical pragma: a portion of code that only thread at a time.
Source of inefficiency: Update to area inside a critical section.
Reduction (<op> :<variable>)

Performance Improvement#1: Too many fork/join lower perform. Inverting loops may help performance if: parallelism is in inner loop. After inversion, the outer loop can be made parallel. Inversion does not significantly lower cache hit rate.
Improvement#2: If loop has too few iterations, for/join overhead is greater than time savings from parallel execution.
#pragma omp parallel for if(n > 500)
Improvement#3: Schedule clause to specify how iteration of a loop should be allocated to threads; static/dynamic inside loop.
Static/Dynamic Scheduling: low/high overhead, imbalance/not
Chunk: is a contiguous range of iteration. Increasing chunk size reduces overhead and may increase cache hit rate. Decreasing chunk size allows finer balancing of workloads.


#schedule (<type>[,<chunk>]). Allowed types:static/ dynamic/ guided/runtime/
Functions for SPMD-style Programming:
Omp_get_thread_num: return ID
Omp_get_num_threads: return total active num of threads
#pragma omp single: only a single process, no matter who.
Nowait clause: a barrier synchronization at end of every parallel for statement.
#pragma omp parallel sections
#pragma omp section
some code here, pack in a section

Characteristic	OpenMP	MPI
Suitable for multiprocessors	Yes	Yes
Suitable for multicomputers	No	Yes
Supports incremental parallelization	Yes	No
Minimal extra code	Yes	No
Explicit control of memory hierarchy	No	Yes


Quinn 18
Why C+MPI+OpenMP faster:
1. Lower communication overhead. 2. more portion of program may be practical to parallelize. 3. May allow more overlap of communications with computations

Algorithms and Sorting
Sorting: $O(n \log n)$ optimal for any sequential sorting algorithm. $O(\log n)$ optimal for parallel.
1. Compare & exchange: 1.P1 send A to P2, P2, compare A with B. 2. both send, both do compare. Notice that different **precision** could conceivably produce different answers.
2. Bubble sort: largest number move to end of list by comparisons. Number of operation: $n(n-1)/2$, $O(n^2)$; odd even
3. merge sort: seq $O(n \log n)$, para $2 \log n$ steps, each step more than 1 operation.
4. Quick sort: seq $O(n \log n)$, choose pivot. May not balance in parallel. Pivot selection is critical

Batcher's parallel sorting: Odd-even merge sort & bitonic mergesort. Both well balanced and have para $O(\log^2 n)$ with n processors
Odd-even merging of two sorted lists
Bitonic Sequence
A monotonic increasing sequence is a sequence of increasing numbers.
A bitonic \rightarrow one increasing and one decreasing.



Feature: If we perform a compare-and-exchange operation on a_i with $a_{i+n/2}$ for all i , where there are n numbers in the sequence, get **TWO** bitonic sequences, where the numbers in one sequence are **all less than the numbers in the other sequence**
Sorting a bitonic sequence



#steps: $n=2^k$, k phases, $\text{steps} = \log(\log n + 1)/2 = O(\log^2 n)$
Summary: odd-even transposition: $O(n)$. Parallel merge: $O(n)$ but unbalanced. Parallel quicksort: $O(n)$, unbalanced, can be $O(n^2)$. Odd-even mergesort and bitonic merge: $O(\log^2 n)$
Sorting on specific network: Algorithms can take advantage of the underlying interconnection network of the parallel computer. Two network structures have received specific attention: **the mesh and hypercube** because parallel computers have been built with these networks.
MPI does provide features for mapping algorithms onto meshes, and **one can always use a mesh or hypercube algorithm** even if

the underlying architecture is not the same.
Shear sort: $\sqrt{n}(\log n + 1)$ for n number, on $\sqrt{n}(\log n)$ mesh
Rank sort: seq $> O(n \log n)$, para $O(\log n)$ with n^2 processors, $O(n)$ with n processors.
Counting sort: a stable sorting algorithm. requires $O(\log n)$ time with $n - 1$ processors. The final sorting stage can be achieved in $O(n/p)$ time with p processors or $O(1)$ with n processors.
Radix sort: Already mentioned parallelized counting sort using prefix sum calculation, which leads to $O(\log n)$ time with $n - 1$ processors and constant b and r .